

ControlPlane

Detailed Business Proposal

Govern the **episode**, not the response - a rupee-denominated risk budget, claim provenance and an action gate across the whole AI task.

Problem

Solution

Users

Business case

Go-to-market

Roadmap

Risks

 **This document**
the detailed business proposal the brief requires

 **Working prototype**
public repo · gateway, console, 10-act demo, eval suite

 **Pitch presentation**
proposal + live prototype, replay-first (offline-safe)

Team Pluoton · August 2026 · every figure measured on the prototype and reproduced by `python -m evals.run`

Round 1 concept selected to advance: a tiered Responsible AI Checker. Round 2 ships it, measures it, and adds the layer no incumbent governs - the episode.

3,488
seeded eval records

26
tests passing

3 + 3
policy packs · jurisdictions

10 acts
offline scripted demo

1 cmd
docker compose up

MEASURED ON THE WORKING PROTOTYPE

41/41
tainted irreversible actions held before execution - 0/47
false holds

96.4-100%
recall per risk category, with 95% confidence intervals

12.6 ms
median added latency per request · 108 req/s per instance

4.8×
ROI floor from remediation alone - before any regulatory term

prototype · console · demo · evals - all in the public repository



CONTENTS

Index - eight sections plus appendix; 02 to 08 are the brief's required proposal spec, in order



Introduction

01 · p.3

Round 1 → Round 2 continuity, the mandate, stated assumptions

6 R1 claims reconciled - every one beaten or made precise



Problem framing

02 · p.4

Detection is commodity; compounding episode risk is ungoverned

Exhibit: compounding-risk curve · 4 incident stats



Solution design

03 · p.5

Risk budget in ₹ · claim taint · action gate - over a policy-driven gateway

Exhibit: 7-step pipeline · tier ladder · 6 decision outcomes



Target users

04 · p.6

CIO + CRO buyer; three governed lanes with measured proof per lane

Exhibit: recall bar chart · 0 false flags · 43/43 abstentions



Business case & impact

05 · p.7

4.8× ROI floor, scenario band shown separately, 0.04% unit economics

Exhibit: ROI chart · assumption stack · 5 unit-economics KPIs



Go-to-market

06 · p.8

Land-expand-run on Accenture's delivery motion; the competitive map

Exhibit: land-expand-run flow · competitive 2×2 matrix



Phased roadmap

07 · p.9

Episode layer live from day zero; phases gate enforcement authority

Exhibit: 18-month gantt · 4 contract-grade milestones



Key risks & mitigations

08 · p.10

Six risks, each countermeasure already shipped - plus stated scope

Exhibit: residual-risk register · 3 verification commands



Appendix

p. 11

Every number in this deck, its source file, and how to reproduce it

Full eval table with CIs · repro commands · assumption register

Brief-required core (sections 02-08, tinted): problem framing · solution design · target users · business case and impact · a phased roadmap · key risks with mitigations - exactly the deliverable specification of the Round 2 problem statement.

We built the prototype first, measured it, and wrote this business case on the results

ROUND 1 PROMISED → ROUND 2 MEASURED

Round 1 claim	Round 2 evidence
Tiered Tier 0-3 checker, inline, real-time	✓ shipped - deterministic → small models → LLM judge → human
< 60 ms added latency (Tier-0 target)	✓ beaten - 12.6 ms p50 / 20.3 ms p95, full lite ensemble
> 90% of injected failures caught	✓ beaten - 96.4-100% recall per category, with 95% CIs
< 3% of model token spend	✓ beaten - 0.04% of model spend under load
“Real-time streaming interception”	made precise: sentence-buffered release, PII masked pre-wire
Bias / responsibility guarantees	✓ hardened - bias stays annotate-only, structurally enforced + regression-tested

Prototype scale: 53 tests · 3,488-record eval set · 10-act demo · React operator console · one-command Docker run

STATED ASSUMPTIONS - BRIEF PARAMETERS, ADAPTED FREELY

3

concurrent AI lanes: customer support · internal copilot · regulated decision-support

2.1M

interactions per year (~40,000/week combined) - the brief's own volume

MIXED

well-governed and loosely governed data sources feed the lanes

API

foundation models consumed via API - input/output-layer inspection only

What changed since Round 1: inline response checking became commodity - leading platforms ship sub-200 ms guardrails. We kept ours, stopped selling it as the innovation, and moved the differentiator to episode-level governance: a ₹ risk budget, claim provenance and an action gate across the whole task.

HOW TO READ - THE SEVEN CORE SECTIONS THAT FOLLOW


03 Problem**04 Solution****05 Users****06 Business case****07 GTM****08 Roadmap****09 Risks**

Each chip is one slide; every quantitative claim on them is annotated with where it is reproduced - the eval suite, the load test, or a stated assumption (full map in the appendix).

Per-response guardrails are commodity - the compounding risk of a whole task is ungoverned

HOW AN EPISODE FAILS WITH EVERY RESPONSE INDIVIDUALLY “SAFE”

Cumulative expected loss across one demo episode (₹ thousand)



Turn 7: HOLD - ₹103,467 vs ₹100,000 budget (measured, demo act 6). No single turn ever crossed a per-response block threshold. Intermediate points illustrative; endpoints measured.

ALREADY SOLVED BY THE MARKET - WE DON'T RESELL IT

- ✓ <200 ms in-VPC response checks
- ✓ EU AI Act / NIST RMF / ISO 42001 policy mapping
- ✓ Hyperscaler guardrails bundled ~free

2026: AGENTS ARE FAILING IN EXACTLY THIS SHAPE

10 / 11

open-source coding & computer-use agents bypassed raw-string shell guards (GuardFall) 

1 email

was enough to poison agent memory - one planted “fact” steered later actions (MemGhost) 

~20%

of organisations report mature AI governance (Deloitte, 2026) 


Art. 14

EU AI Act demands effective human oversight - for agents, only dischargeable at episode level 

“Multi-turn conversations and AI agents that take actions introduce compounding risk, where one questionable output can shape several downstream decisions.”

- Round 2 brief, Problem Track 1: the gap, named by the examiner

Three mechanisms govern the episode - on a gateway where policy is data, not code



OpenAI-compatible proxy · per-use-case latency budgets · PII masked before it reaches the wire · every decision hash-chained and HMAC-signed

TIER 0 · deterministic - regex, checksum, taint · ~0 ms

TIER 1 · lexical & small models - in budget

TIER 2 · LLM judge - >96% need no LLM call

TIER 3 · human - override queue, 2-person rule

1

Episode risk budget - in rupees

Every passed-but-uncertain output debits a hazard-based expected-loss meter, priced from a severity table the client's risk office signs. The EPISODE escalates when cumulative risk crosses budget.

demo: HOLD at turn 7 - ₹103,467 vs ₹100,000

2

Claim provenance & taint

Numbers, dates, names, IDs are canonicalised - “eighty-five thousand” = 85000 = ₹85,000. A value born ungrounded in model output stays tainted across every later turn. Deterministic, milliseconds, no model call.

words → digits cannot launder a fabrication

3

Action gate on reversibility

An irreversible tool call - pay, send, delete, submit - needs a taint-clear episode, not just clean arguments. Held BEFORE execution, evidence chain attached, two-person override.

measured: 41/41 tainted held · 0/47 false holds

ONE CALIBRATED DECISION PER RESPONSE - SIX OUTCOMES

PASS

ANNOTATE

REPAIR

ESCALATE

BLOCK

HOLD ACTION

REPAIR is deterministic (mask / correct / cite-strip) - the “edit” tier that latency pressure never cuts.

GOVERNANCE AS CONFIGURATION

Three signed policy packs + IN / EU / US jurisdiction overlays, hot-reloaded with anti-rollback. Detection stays pluggable commodity - bring any vendor; the episode layer above it is the product.

customer_support · fail-open

internal_copilot · 1.5 s

decision_support · fail-closed

One platform, three governed lanes - each lane a policy pack with its own physics

Economic buyer: CIO / Head of AI Platform, with the CRO as co-sponsor whose regulatory pressure creates urgency. Engineering, FinOps and operations are stakeholders to neutralise - latency, cost, alert fatigue - not signatures to collect. Daily users: risk-ops analysts on the console + the L1/L2 triage bench.

Customer support assistant



customer-facing · real-time

- ✓ 150 ms detector budget
- ✓ fail-open on checker fault
- ✓ PII masked mid-stream
- ✓ REPAIR over BLOCK - protect CX

privacy 100% · 0 false flags

Internal knowledge copilot



employee-facing · near-real-time

- ✓ 1.5 s budget - fuller ensemble
- ✓ source-trust tiers debit differently
- ✓ annotate-first posture
- ✓ abstains instead of inventing

grounding 96.4% · abstention 43/43

Decision-support tool




regulated · action-taking

- ✓ gate mode · fail-closed
- ✓ actions held on episode taint
- ✓ two-person override
- ✓ IN/EU/US packs flip outcomes live

41/41 held · 0/47 false holds

MEASURED RECALL BY RISK CATEGORY - % OF INJECTED FAILURES CAUGHT



0


false flags across 2,381 benign test records - over-flagging is the adoption killer; we measured against it

43/43

correct abstentions - INSUFFICIENT_EVIDENCE instead of an invented answer, per the brief's "no ground truth" complexity

₹1.2 Cr avoided loss against ₹25L all-in cost - a 4.8× floor before any regulatory term

FLOOR CASE, PER CLIENT PER YEAR (₹ CR)



ROI floor $\approx 4.8\times$ remediation only, zero regulatory term

EVERY INPUT ON THE SLIDE

Volume	2.1M interactions/yr (brief parameter)
Material-failure rate	1.0% - assumed, then measured in shadow
Catch rate	95% - measured on the benchmark
Cost per failure	₹600 remediation + concession
All-in cost side	3 × ₹5L platform + ₹10L managed service, incl. triage bench (8,400 esc/yr × 10 min × ₹500/hr $\approx$ ₹7L)

Scenario band - separate, never blended: probability-weighted DPDP / disclosure / goodwill exposure adds ₹30L-₹1.5Cr. We never anchor on statutory ceilings.

WHY THIS SURVIVES A CFO

- ✓ Attack the assumption, not the method - all inputs visible
- ✓ Failure rate isn't ours to claim: shadow mode measures it on the client's traffic
- ✓ Severity table isn't ours to invent: the client's risk office signs it
- ✓ Triage labour sits inside the denominator - ROI is net of humans
- ✓ The budget meter IS the ROI dashboard - same ₹ machinery, live
- ✓ Volume is the brief's own parameter - no invented market size

UNIT ECONOMICS - WHY THE MARGIN STORY HOLDS AT SCALE

0.04% 

assurance compute vs model spend (load test)

>96% 

of traffic checked with no LLM call at all

12.6 ms 

p50 added latency · 20.3 ms p95

108 req/s 

per stateless instance at saturation · 0 errors

0.065 

expected calibration error - scores you can price risk on

Land with an audit, price the governed use case, run the service - on Accenture's own motion

LAND - Assurance Assessment

4-6 weeks · fixed fee · AUDIT mode over exported logs - zero production footprint, no InfoSec review blocking week 1. Deliverables: episode risk report + client-signed severity table.



EXPAND - In-tenancy deployment

priced per governed use case / agent class per month, tiered by risk class & jurisdictions. 0.04% COGS is the margin story, never the price anchor.




RUN - Managed assurance

L1/L2/L3 triage bench · quarterly policy-pack updates · drift recalibration vs a frozen canary set. Human review IS the recurring line - priced in.



COMPETITIVE MAP - THE TOP-RIGHT QUADRANT IS EMPTY



THE MOAT IS THE UNIT OF ACCOUNT



Bring-your-own-detector is a feature: commoditised detection makes the neutral episode & evidence layer above it MORE valuable. A fast-follow by incumbents requires abandoning per-response pricing and architecture - their revenue model is the switching cost.

unit of account

signed severity method

delivery install base


WHY ACCENTURE WINS WITH THIS



It productises what the Responsible AI practice already does manually. Thousands of GenAI delivery engagements are the install base; the assessment, the signed severity methodology and the managed bench are services only a delivery firm can sell at scale.

Worked example at the brief's volume: ≈ ₹25L/yr per client vs ₹1.2Cr+ modelled avoided loss (section 06).

The episode layer runs from day zero - phases only expand its authority to act



Six risks, six shipped countermeasures - plus scope stated before you ask

False-positive fatigue → users bypass the checker

MED

Shadow-first · operating point is the client's dial on a measured sweep · 0 false flags / 0/47 false holds measured · overturn-rate watch

Latency regression breaks customer experience

LOW

Per-detector timeouts under each pack's budget · degraded coverage debits MORE risk · 12.6 ms p50 measured · stateless scale-out

The checker itself fails or is attacked

MED

Fail-open (chat) / fail-closed (regulated) per pack · Tier 0 never sheds · ingress gate blocks poisoned documents · GuardFail-informed red-teaming

Detector drift / model or vendor change

MED

Weekly recalibration vs a frozen canary set · model fingerprint on every decision · pluggable adapters - swap vendors, keep governance

Evidence ledger becomes a privacy liability

LOW

Keyed-HMAC digests, no raw text at rest · quarantined feedback store · retention policy per jurisdiction pack

Incumbents fast-follow the episode concept

MED

Moat = unit of account + signed severity methodology + Accenture install base - following requires abandoning per-response pricing

STATED SCOPE - SAID BEFORE YOU ASK



Eval rates come from a 3,488-record seeded synthetic benchmark - labels true by construction; real-traffic rates are established in the shadow phase (that is what it is for). Lite profile = deterministic + lexical detectors; NLI and LLM-judge adapters are integrated but optional - values under 10 (a '3-year' vs '1-year' claim) are their job, not the lexical path's. A blind hold-out slot, authored by whoever didn't build the detectors, is scored separately.

THE ASK - VERIFY IN 10 MIN



```
docker compose up --build
```

gateway + operator console on :8080

```
python -m demo.run_demo
```

the 10-act scripted demo, fully offline

```
python -m evals.run
```

reproduce every number in this proposal

Every number in this deck, its source file, and the command that reproduces it

MEASURED RESULTS - TEST SPLIT, 2,381 OF 3,488 SEEDED RECORDS



Risk category	Recall	95% CI	False flags
Grounding (unsupported claim)	96.4% (402/417)	[94.2, 97.8]	0 / 1,964
Privacy (PII)	100% (132/132)	[97.2, 100]	0 / 2,249
Toxicity / harmful language	100% (96/96)	[96.2, 100]	0 / 2,285
Prompt injection (incl. indirect)	100% (68/68)	[94.7, 100]	0 / 2,313
Cost anomalies	100% (26/26)	[87.1, 100]	0 / 2,355

Action gate 41/41 tainted held · 0/47 false holds | abstention 43/43 | ECE 0.065 | labels are true by construction (failures injected on purpose)

WHERE EACH DECK NUMBER LIVES



96.4-100% recall · CIs · 0 false flags · ECE	<code>evals/out/results.json</code>
Recall vs FP threshold curves (operating point)	<code>evals/out/operating_sweep.json</code>
12.6 ms p50 · 20.3 ms p95 · 108 req/s · 0.04% spend	<code>evals/out/load_test.json</code>
₹103,467 vs ₹100,000 turn-7 escalation	<code>demo act 6 (deterministic, offline)</code>
4.8× ROI floor · ₹25L cost · scenario band	<code>stated assumptions (section 05)</code>

REPRODUCE EVERYTHING



```
python -m pytest tests/
```

53 tests on the mechanisms

```
python -m evals.generate
```

3,488-record dataset, byte-identical

```
python -m evals.run
```

recall, CIs, calibration, sweep

```
python -m demo.run_demo
```

10-act scripted demo, offline

ASSUMPTION REGISTER



- ✓ Volume 2.1M interactions/yr - the brief's own parameter
- ✓ Material-failure rate 1.0% - measured in shadow phase
- ✓ ₹600 per failure reaching a user (remediation + concession)
- ✓ Severity table signed by the client's risk office at land
- ✓ Model cost ₹0.04/1K tokens · judge call ₹0.05 (config.py)
- ✓ Severities and budgets are policy-pack data, not code