

Every data point in this proposal and exactly where it comes from

MEASURED - FILE IN THE REPOSITORY



Grounding recall 96.4% (402/417) + 95% CI	evals/out/results.json
Privacy / toxicity / injection / cost recall 100%	evals/out/results.json
0 false flags per category (1,964-2,355 benign)	evals/out/results.json
Action gate 41/41 held · 0/47 false holds	evals/out/results.json
Abstention 43/43 · ECE 0.065	evals/out/results.json
Recall vs FP-rate curves by threshold	out/operating_sweep.json
Per-detector probability calibration	evals/out/calibration.json
Coverage weights (risk-weighted recall)	out/coverage_weights.json
12.6 ms p50 · 20.3 ms p95 added latency	evals/out/load_test.json
108 req/s saturation · 0.04% of model spend	evals/out/load_test.json
Dataset: 3,488 records, 2,381 test, ~28% injected	evals/generate.py (seed 20260823)
Turn-7 escalation ₹103,467 vs ₹100,000	demo/run_demo.py - act 6
>96% of checks with no LLM call	demo act 10 + results.json
53 passing tests	tests/ (python -m pytest)

STATED ASSUMPTIONS - WHERE DECLARED



Volume 40,000/week = 2.1M interactions/yr	Round 2 brief, reference parameters
Three lanes: support, copilot, decision-support	Round 2 brief, reference parameters
Mixed well-/loosely-governed sources	Round 2 brief, reference parameters
Material-failure rate 1.0% (measured in shadow)	proposal section 05 - assumption
Catch rate 95% (benchmark-derived, rounded)	proposal section 05 - assumption
₹600 remediation + concession per failure	proposal section 05 - assumption
All-in cost ≈ ₹25L/yr incl. ₹7L triage bench	proposal section 05 - worked example
ROI floor 4.8× · band ₹30L-₹1.5Cr	arithmetic over the rows above
Episode budget ₹100,000 · severity table	decision_support.yaml
Latency budgets 150 ms / 1.5 s / gate mode	policies/*.yaml (3 packs)
Model cost ₹0.04/1K tokens · judge ₹0.05	controlplane/config.py
Severities signed by client risk office at land	GTM design, section 06

EXTERNAL & PUBLIC SOURCES



"Compounding risk" quote & brief requirements	AIC 2026 Round 2 problem statement, Track 1
10/11 agents bypassed shell guards	GuardFall - 2026 agent-security research (public)
Agent memory poisoned by one email	MemGhost - 2026 agent-security research (public)
Review modes steered by hostile repo content	Friendly Fire - 2026 agent-security research (public)
~20% of organisations report mature AI governance	Deloitte AI governance survey, 2026
Human-oversight duty for high-risk AI	EU AI Act, Article 14
Hallucination failure archetypes in eval data	HaluEval · RAGTruth (public benchmarks)
Jailbreak / injection archetypes in eval data	JailbreakBench (public benchmark)
PII patterns in eval data	Microsoft Presidio (open source)
Sub-200 ms guardrails are commodity; session observability	Fiddler · Galileo · Arthur public positioning, 2026
Bundled content guardrails	AWS Bedrock / Azure AI public docs, 2026
Gateway guardrail hooks	LiteLLM · Portkey public docs, 2026
Third-party attribution	repository NOTICE file (Apache-2.0)

Reproduce the measured column from a fresh clone: `python -m pytest tests/` → `python -m evals.generate` && `python -m evals.run` → `python -m scripts.load_test`
 - the dataset is seeded, so every number regenerates byte-identically.